

AI ALLIANCE

A global non-profit organization building, supporting and advocating for open research and open source AI.

The AI Alliance is working to *scale intelligence outward* so every person and business has the capability, control and customization they want.



Our Community

205

Organizations*

From startups to enterprises



29

Countries

Truly global reach



50+

Events Annually

Driving engagement



100+

Projects

Building the foundation



25K+

Community Members



Our Program

Program Pillars

Agent Hub blueprints, tools and guidance for building domain-specific agents.

Open Data and Models to enable frontier sovereign AI that you own and control.

AI Safety, Governance, and Education to ensure AI is useful, accessible and trusted.

AI Alliance Community
501(c)(3)

Open Source AI Innovation Association convene, activate, and advocate for open standards, practices and policies.

AI Alliance Association
501(c)(6)

Example Projects: Technology



Example Projects: Education, Advocacy and Policy



Explore our projects: github.com/The-AI-Alliance

Learn More about AI Alliance



thealliance.ai



github.com/the-AI-Alliance

Project Tapestry

A consortium approach to training frontier foundation models and sovereign derivatives.



TAPESTRY
NATIONAL - CULTURAL - INDUSTRIAL

AI ALLIANCE

GitHub: <https://github.com/The-AI-Alliance/tapestry>

Website: <https://thealliance.ai/projects/tapestry>

Frontier Sovereign Models

Today, models with frontier performance are built by only a few organizations, and who determine their content, capability, and accessibility.

Yet, only a fraction of the world's data, compute infrastructure, and technical community are involved in frontier model development.

What if we could build a new **global foundation model** with this broader set of knowledge, resources, people and organizations working together?

How do we do this?

Project Tapestry will bring together data, compute and talent from a global consortium of partners to build the next frontier of foundation models trained on a larger and more diverse corpus than ever before.

Tapestry will enable sovereign AI by ensuring ownership of data and compute remains with partners, and partners can continue to train sovereign derivatives of the consortium-trained base model that they own.

The Tapestry consortium led by a strong core operating entity will create **a sustainable operating model with a structural advantage that no single organization or centralized lab can easily replicate.**

AI ALLIANCE

Project Tapestry

GitHub: <https://github.com/The-AI-Alliance/tapestry>

Website: <https://thealliance.ai/projects/tapestry>

Performance is the Path to Sovereignty

PROJECT TAPESTRY GOALS

FRONTIER PERFORMANCE

Tapestry's ambition is to create a global foundation model system that is more performant than any other model while enabling sovereign derivatives that are owned and controlled by partners.

SOVEREIGNTY

A person, organization, or nation can own and control their AI, use it how they see fit, and retain the value they create with it, ensuring independence and agency.

STRUCTURAL ADVANTAGE

Consortium-managed access to data used for training creates a structural advantage. Distributed ownership, sustainable funding model and mission-locked governance ensure trust.

AI ALLIANCE


Project Tapestry

GitHub: <https://github.com/The-AI-Alliance/tapestry>

Website: <https://thealliance.ai/projects/tapestry>


Project Tapestry

Consortium-developed open platform and frontier model base.
Partners own most data, sovereign derivatives, and commercial deployment.



Consortium Training: Shared Core, Sovereign Output

The base model improves through shared training, while partners retain control over their data, compute, and sovereign derivatives.



First Technical Milestones

Specific early goals to validate hypotheses and de-risk scale up.

Cultural alignment

Release and publish the first model re-aligned to 2+ cultures without core performance loss.

Distributed Training

Demonstrate and open source the first multi-node weight update and aggregation framework. Run experiments to understand scale up requirements.

Training Data Catalog

Create a registry of metadata associated with candidate data sources potentially available to consortium for training.

📄 .gitignore	feat: add federated training and TVA docs	2 weeks ago	🔗 Custom properties
📄 AGENTS.md	feat: replace federated demo with consortium training PoC	2 weeks ago	★ 15 stars
📄 GITHUB_PAGES.md	Added warning about Ruby 4	3 weeks ago	👁 3 watching
📄 Gemfile	Initial commit	3 weeks ago	🍴 1 fork
📄 LICENSE.Apache-2.0	Initial commit	3 weeks ago	🔍 Audit log
📄 LICENSE.CC-BY-4.0	Initial commit	3 weeks ago	📄 Report repository
📄 LICENSE.CDLA-2.0	Initial commit	3 weeks ago	Releases
📄 Makefile	typos	19 hours ago	No releases published Create a new release
			Packages


GitHub: <https://github.com/The-AI-Alliance/tapestry>

Website: <https://thealliance.ai/projects/tapestry>

Data Catalog: <https://thealliance.ai/projects/tapestry/training-data-proposals>

Long Term Roadmap

Progression between phases is gated by technical proof points, partner commitments, compute availability, data readiness, and funding thresholds.



Momentum is building quickly.

↕ You reposted



Yann LeCun 
@ylecun

X.com

The salvation is Project Tapestry
thealliance.ai/projects/tapes...



Daniel Jeffries ✓ @Dan_Jeffries1 · 1d

I don't think people understand just how bad it will be if an American open source champion doesn't emerge soon and the big labs succeed in creating modern East Indian companies and ban open models on moronic national security grounds.

1:46 PM · 5/17/26 · **53K** Views

31

↻ 28

♥ 205

77



AI Alliance Advances Project Tapestry as G7 Puts AI Sovereignty at Center Stage

AI ALLIANCE

NEWS PROVIDED BY
AI Alliance →
Jun 19, 2026, 12:09 ET

SHARE THIS ARTICLE



India's BharatGen commits to anchor India's role in Project Tapestry as part of a growing coalition of organizations to develop more capable frontier AI that enhances sovereignty and opportunity for people, nations and industry.

PARIS, June 19, 2026 /PRNewswire/ -- In a week when AI sovereignty dominated discussions from the G7 Summit in Évian to VivaTech in Paris, the AI Alliance announced that India's BharatGen will anchor India's initial role in Project Tapestry, a global open consortium to build frontier-capability AI through distributed model development while allowing participating nations and institutions to retain control over their own data, models, and deployment.

The commitment connects India's self-reliant AI ambitions with Tapestry's collaborative model for building frontier AI while preserving national control over data, models, and deployment.

"India has long believed that science advances fastest when it is open and shared. Project Tapestry, which lets nations build frontier AI together while retaining sovereignty over their data and models, reflects the collaborative, self-reliant innovation India is advancing."

— Professor Ajay Sood, Principal Scientific Adviser to the Government of India



Project Tapestry



thealliance.ai/projects/tapestry



github.com/The-AI-Alliance/tapestry




AI ALLIANCE

— FRONTIER · SOVEREIGN AI

Project Tapestry

A consortium approach to training frontier foundation models and sovereign derivatives.

YANN LECUN · CHIEF SCIENCE ADVISOR

[GET INVOLVED](#)[Alliance ↗](#)

Operating Model

Design principle: decentralized contribution, centralized technical integration, globally trusted governance.

Consortium Governance.

Responsible for overall strategy and roadmap and major role and responsibility decisions.

Partners (many)

- Maintains teams and resources.
- Clear, regulation-compliant contract commitments to the core entity.
- Owns their data and compute.
- Owns specialized evaluation.
- Responsible for local training and selective sharing back to core.
- Responsible for building and updating sovereign derivatives.
- Responsible for developer & partner ecosystem for sovereign derivatives.
- Independent commercialization.

Local investment as part of consortium commitment.

Sovereign model derivatives.



Base model development.



Core Operator

- Core team and resources.
- Board: partners and independents.
- Funders take part in governance.
- Curates distributed data catalog.
- Manages core compute and some data, coordinates global network.
- Responsible for core training & aggregation of distributed training.
- Responsible for building and updating global base models.
- Responsible for developer & partner ecosystem for base model.

Central investment sustained to support core work.

Training Data Catalog

The data we bring together may become one of Tapestry's most important assets and the source of our differentiation and structural defensibility. As a starting point we seek to catalog potential sources of training data from partners.

What can be contributed?

- public or private data sets
- domain/industry data
- language and cultural resources
- government/public records where lawful
- evaluation datasets and environments
- preference and feedback data

What must be described?

- provenance and ownership
- license and use restrictions
- privacy / PII status
- quality and AI-readiness
- geography and jurisdiction
- permitted use: pre/post-training, evals, synthetic gen


Benefits to Contributors

- recognition and contribution accounting
- access to improved base models
- right to create local derivatives
- control over restricted data use
- path to commercial and sovereign ROI
- governance input

Propose Data Sets Here: <https://thealliance.ai/projects/tapestry/training-data-proposals>

Consortium Training: A New Approach


Data from partners can be contributed privately to the the core training corpus and used only for base model training. Sensitive data can be held back at the partner node and used locally to update weights, with only the weights shared.



Independent Tuning and Aligning By Partners

Any partner can create a derivative model of the base that they fully own and control using the same infrastructure.

Fine tuning (post-training) and
Reinforcement Learning (RL) for
alignment



Partner


- Continued pre-training, SFT, RL, etc.
- Use private or other locally-relevant data and environments.
- No updated model weights shared with core.

The same “stack” is
used for local tuning
with private and other
data sets.

The Tapestry stack supports tuning the base models to build custom, local models that use private and other data. Examples:

- A national model aligned to local laws and culture.
- An industrial model tuned on manufacturing procedures.
- A healthcare or government model trained in a secure private environment.

LLM Training Process



LLM Training Dataset Types

raw unprocessed partner sources (corpora, records, crawls)	prepared cleaned, structured text + metadata	pre-training tokenized training mixes	SFT (supervised fine tuning) instruction / chat data
preference chosen vs. rejected responses	reward-verifiable problems with checkable rewards	validation held-out splits per stage	eval benchmarks and environments


Tapestry: Metadata-First Contribution

Partners describe contributions by type + metadata from the catalog — not by shipping raw corpora. Data stays local; the consortium negotiates gaps (“more prepared in language X”, “node Y’s eval for cultural alignment”).

The catalog tells us what exists and how it must be described. These types tell us how the consortium names it, gates quality, and matches contributor rights to what was actually contributed.

Tapestry: Training while Adhering to Data Requirements

Categories of Data Sovereignty Supported by Tapestry



Project Tapestry



thealliance.ai/projects/tapestry



github.com/The-AI-Alliance/tapestry

Our Community

205

Organizations*

From startups to enterprises

29

Countries

Truly global reach

50+

Events Annually

Driving engagement

100+

Projects

Building the foundation

25K+

Community Members



Our Program

Program Pillars

Agent Hub blueprints, tools and guidance for building domain-specific agents.

Open Data and Models to enable frontier sovereign AI that you own and control.


AI Safety, Governance, and Education to ensure AI is useful, accessible and trusted.

**AI Alliance
Community**
501(c)(3)


Open Source AI Innovation Association convene, activate, and advocate for open standards, practices and policies.

**AI Alliance
Association**
501(c)(6)

Example Projects: Technology



Example Projects: Education, Advocacy and Policy



Explore our projects: github.com/The-AI-Alliance

Learn More



thealliance.ai



github.com/the-AI-Alliance

Projects



Deep Research Agent for Finance

Goal: Provide a reference implementation for building production-quality agents with sophisticated reasoning that access data and tools only through MCP servers.

Agent determines which financial information sources to query, aggregates and analyzes results, and provides an analyst's report.

Technology Partners: Currently working with Kasikorn, S&P, and BrightQuery to improve and broaden research capabilities.

The screenshot displays the 'Deep Finance Research' interface. It features a 'Task Queue' on the left with a tree view of 'Completed Steps' (Segments and KPIs Collection, Compile segments and KPIs including DAW, Validation and Documentation, Perform checks including sum checks, rat, Document all sources and references used) and an 'Active Step' (Output Compilation, Pack all structured financial data into ...). Below the task queue is a 'Budget' table showing resource usage.

Resource	Used	Limit	Usage %
Tokens	102,280	100,000	102.3%
Cost	\$0.102	\$1.00	10.2%
Time	6.0 min	10 min	67.7%

Below the budget table, a 'Status' section shows the 'Objective' (You are a meticulous financial analyst specializing in tech company financial research. Your role is...) and 'Iteration: 8/25'. A 'Policy Engine' section shows 'Consecutive Failures: 0/3', 'Total Successes: 7', 'Total Failures: 0', and 'Failure Rate: 0.0%'. An 'Agent Cache' section is also visible.

At the bottom, a 'Tokens' section shows 'usage' (113,052 tokens | \$0.0294) and a 'Finished' status for 'FinalSynthesizer' with an 'Elapsed Time' of 00:00:32. A 'Chatting' section shows 'ExcelWriter' with 'o4-mini turn 1'.



lastmile AI

IBM



Retail Analytics Agent

Natural Language Automation

Architecture for retail analytics with structured transaction data, supporting most off-the-shelf retail management frameworks including SAP Retail.

Graphical presentation rivals static business intelligence dashboards.

Roadmap: Broader and local model support, open source UI.



Semiconductor Manufacturing Agent

Goal: Develop a deep semiconductor domain expert agent that can advise engineers on etch development, root cause defect analysis, and debugging.



Domain-Tuned Model

Custom evaluation with expert data co-developed with workflow



Deterministic Grounding

Outputs grounded in library of known-good processes



Dana Framework

Developed within Dana agent framework and runtime

SiO₂ Etching Advisor
powered by SemKong

PlasmaPro 100 Cobra ICP RIE Etch



REQUIREMENTS & SPECIFICATIONS

I am trying to etch 2 µm of PECVD SiO₂ using a -4 µm PR mask to create a pattern of 20 * 60 µm. Recommend me 2 recipes.

[Get Recipe Advice](#)

RECIPES

Recipe 1	Recipe 2
Gases and Flow Rates: CHF₃: 40 sccm	Gases and Flow Rates: CHF₃: 50 sccm

QUALITY & SAFETY ADVICE

- Always follow safety procedures when handling gases and operating plasma etching systems.
- Confirm compatibility of recipes with your specific equipment to avoid damage or contamination.

TEL


AITOMATIC



Semiont Knowledge Platform

Semiont enables humans and agents to co-create shared knowledge through an AI-native Wiki and Annotator.

- Graph of Documents and W3C Web Annotations
- High accuracy context retrieval
- Integrates via SDK & MCP
- Locally owned knowledge bases deployable where you want



Dana: Sovereign Agent Framework

A new **sovereign agent development and deployment framework** with native support for agentic workflows, memory, grounding, and concurrency.

- Integrates domain-specific knowledge from human expertise
- Designed to run locally or in cloud
- AI models + symbolic grounding for reliable outputs



Open Trusted Data Initiative

Metadata Specification

For provenance, lineage, and utility

Catalog

Existing open data sets based on trust score

Curation and Creation

New data sets needed for agentic use cases

Project
Leadership



NextGem Data Catalog



Minimum Score

0

0

Enter use case...

Enter language...

Enter license...

Search Catalog **Reset**

[Join Our Initiative](#) [Browse the Datasets](#) [Contribute a New Dataset](#)

Datasets Catalog / Modality Datasets

Datasets for Different Modalities

Modalities include text, vision, different widely-applicable concepts, like data formats, how the data was collected or transformed from other data (e.g., see text-to-..., etc.), and general usage guidance like data intended for pretraining, reinforcement-learning, chain of thought, etc.

Keywords

3D Agents Alignment Arrow Anxv Audio Benchmark Classification Chain Of Thought

Chat Crowd Sourced CSV Embeddings Evaluation Fine Tuning Generated Data

Feature Extraction Graph Handwritten Image Instruction Following LLM JSON Monolingual

Multi Lingual Multimodal Multiple Choice Named Entity Recognition News NLP Planning

Pretraining Problem Solving Prompt Question Answering RAG Reasoning Regression

Reinforcement Learning Safety Search Security Sentence Similarity Sentence Transformers

Sentiment Analysis Speech Summarization Tabular Retrieval Text To ... To Text Translation

Tutorial Unlearning Video Vision Wikipedia

Synthetic Data for Open Frontier AI

- [illegible]



Semikong: Semiconductor Foundation Model

An open foundation model for **semiconductor** applications — trained and evaluated with data collected from industry expert SMEs, co-designed to enable trusted agentic assistants.

Expert Training

Data collected from industry SMEs

Agent-Ready

Co-designed for process engineer assistants

Fully Open

Model weights, inference code, and data sets



SEMİKONG
Empowering Semiconductor Innovation
Meet Semikong, the World's First Semiconductor Industry-Specific Large Language Model
[Download Semikong 7B](#)

SEMİKONG IS A COLLABORATIVE OPEN-SOURCE PROJECT BORN FROM THE AI ALLIANCE.
SPECIAL THANKS TO OUR SUPPORTERS FOR HELPING MAKE IT HAPPEN.

SUPPORTERS AUTOMATIC TEL Software Accelerator AI Alliance

Designed for Semiconductor Unique Challenges
TRAINED WITH SEMICONDUCTOR KNOWLEDGE

-  Trained on a comprehensive semiconductor-related text corpus
-  Novel pre-training approach leveraging domain-specific knowledge
-  Superior performance compared to general-purpose LLMs on industry-relevant benchmarks
-  Serves as a valuable foundation for companies to build proprietary models tailored to their needs

